

Question 1

1. (15 points) Consider a dataset with three data points in \mathbb{R}^2 :

$$X = \begin{bmatrix} 0 & 0 \\ 0 & -1 \\ -2 & 0 \end{bmatrix} \quad y = \begin{bmatrix} -1 \\ -1 \\ +1 \end{bmatrix}$$

Manually solve the optimal hyperplane optimization problem (below) to get the optimal hyperplane (b^*, w^*) and its margin.

Optimization Problem:

$$\min_{b, w} \frac{1}{2} w^T w$$

subject to: $y_n(w^T x_n + b) \geq 1 \quad (n = 1, \dots, N)$

We rewrite X and y as a system of inequalities

$$\bullet -b \geq 1$$

$$\bullet +w_2 - b \geq 1$$

$$\bullet -2w_1 + b \geq 1$$

$$\Rightarrow -w_1 + \frac{1}{2}w_2 \geq 1$$

$$\Rightarrow -w_1 \geq 1 \quad = w_1 \leq -1$$

$$\Rightarrow w_2 \geq 0$$

Using our minimization condition

$$\min_{b, w} \frac{1}{2} w^T w = \min_{b, w} \frac{1}{2} (w_1 \ w_2) \begin{pmatrix} w_1 \\ w_2 \end{pmatrix} = \min_{b, w} \frac{1}{2} (w_1^2 + w_2^2)$$

By inspection, our minimizing values would be

$$w_1 = -1$$

$$w_2 = 0$$

$$b = -1$$

The margin is simply:

$$\frac{1}{\|\vec{w}\|} = \frac{1}{1} = \boxed{1}$$

$$\Rightarrow w = \begin{pmatrix} -1 \\ 0 \end{pmatrix}, b = -1$$

Question 2

2. (15 points) Consider the following proof concerning the matrix Q from the linear hard-margin SVM algorithm.

Proof: Let $\mathbf{u} = [x_0, \mathbf{x}_u^T]^T \in \mathbb{R}^{d+1}$ and $\mathbf{x}_u \in \mathbb{R}^d$. Since we have $Q = \begin{bmatrix} 0 & 0_d^T \\ 0_d & I_d \end{bmatrix}$, then

$$\mathbf{u}^T Q \mathbf{u} = [x_0, \mathbf{x}_u^T] \begin{bmatrix} 0 & 0_d^T \\ 0_d & I_d \end{bmatrix} \begin{bmatrix} x_0 \\ \mathbf{x}_u \end{bmatrix} = \mathbf{x}_u^T \mathbf{x}_u \geq 0$$

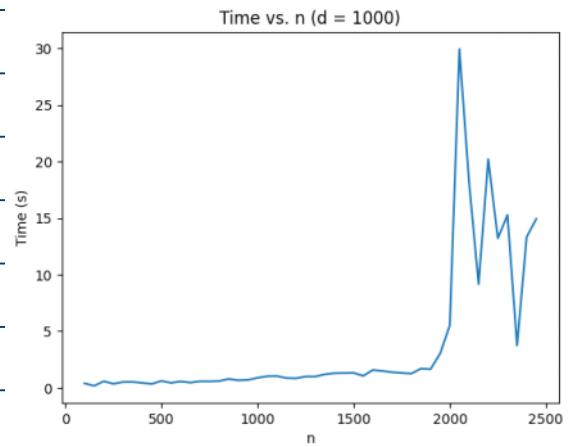

This is the case for arbitrary \mathbf{u} .

- (a) What property of matrix Q does this prove?
- (b) What does this property of Q mean for the standard QP problem?
- (c) What is the usefulness of (b) in terms of finding a solution for the QP problem?

- a) By definition it makes Q a positive semi-definite matrix
- b) It makes the QP problem a convex problem
- c) It helps us find the optimal solution easier being able to utilize properties of convex problems to use certain well known methods. Specifically QP -solvers can solve it in $O((N+d)^3)$

Question 3

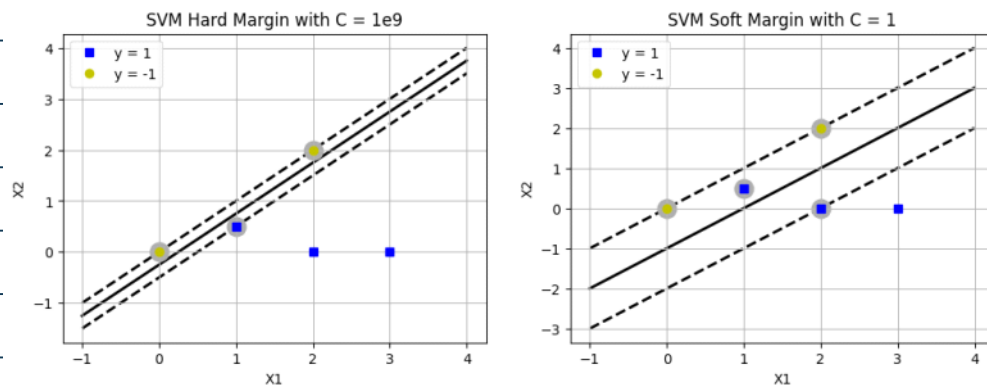
Part b:



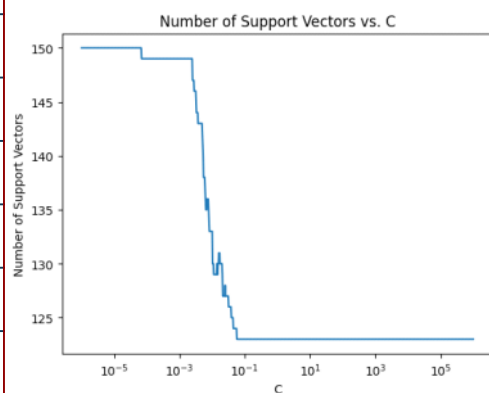
Our algorithm is not guaranteed to be efficient given any size of X . As we go to higher numbers of features or datapoints we get an algorithm that grows more than linearly. The specific big-oh notation is not obvious from the tests we have done above and would include more analysis into the calculation itself.

Question 4

Part a: In this part I add a point to the test data before that would be on the wrong side of the margin. Here I train both hard and soft margin SVM to see the difference. The results is that the hard margin changes a lot to make it still perfectly separable while the soft margin keeps the same margin as before and just allows the point to be misclassified.



Part b:



Here is a plot of the number of support vectors versus C. Notice that the x-axis is in log scale. We see behavior as we would expect where at very low C we get the maximum number of support vectors and as we increase C it decreases

Here I experiment with different kernels and extremes of C values for each

Linear kernel, C = 1, Training Accuracy = 1.0
Linear kernel, C = 1, Testing Accuracy = 0.8698245614035087
Linear kernel, C = 1e-6, Training Accuracy = 0.16
Linear kernel, C = 1e-6, Testing Accuracy = 0.0968421052631579
Linear kernel, C = 1e6, Training Accuracy = 1.0
Linear kernel, C = 1e6, Testing Accuracy = 0.8698245614035087
Polynomial kernel, C = 1, Training Accuracy = 0.9866666666666667
Polynomial kernel, C = 1, Testing Accuracy = 0.8080701754385965
Polynomial kernel, C = 1e-6, Training Accuracy = 0.16
Polynomial kernel, C = 1e-6, Testing Accuracy = 0.0968421052631579
Polynomial kernel, C = 1e6, Training Accuracy = 1.0
Polynomial kernel, C = 1e6, Testing Accuracy = 0.8059649122807018
RBF kernel, C = 1, Training Accuracy = 0.9933333333333333
RBF kernel, C = 1, Testing Accuracy = 0.871578947368421
RBF kernel, C = 1e-6, Training Accuracy = 0.16
RBF kernel, C = 1e-6, Testing Accuracy = 0.0968421052631579
RBF kernel, C = 1e6, Training Accuracy = 1.0
RBF kernel, C = 1e6, Testing Accuracy = 0.8887719298245614

Our results are interesting. For linear, polynomial, and RBF we get bad results for a low C values and better for higher. RBF with a high C value performed best on the testing set.